

Fused Codebook Dequantization for LLM Decode, Cross-GPU Speedups and a Bandwidth Law

Azziz Mimouni

github.com/neuralboot/trapetum

July 2026

Abstract

Autoregressive LLM decode at batch size one is memory bandwidth bound: producing each token requires streaming the entire weight matrix from device memory, while the arithmetic units sit nearly idle. We study a fused 4-bit codebook GEMV kernel that folds the index to weight lookup into the matrix vector product, so the full precision weight matrix is never materialized and only one quarter of the weight bytes are read. Measured against a live cuBLAS half precision GEMV on the same shapes, fixed seed, and verified for numerical correctness, the kernel reaches **2.20x** faster decode on an RTX 4090, **2.34x** on an A40, and **parity (0.99x)** on an H100. The speedup tracks how bandwidth limited the device is: on bandwidth rich parts the half precision GEMV is already near roofline and reading fewer bytes only saves memory. Integrated the way serving engines run, with the decode step captured as a CUDA graph, this realizes a **2.0x** (Llama-2 7B) to **2.45x** (13B) end-to-end decode speedup at roughly a third of the memory. We pair this with a fair single harness benchmark of quantization methods (fp16, AWQ, AQLM, Llama-2 7B and 70B, full wikitext-2 perplexity), where a 2-bit 70B model fits a single 24 GB consumer GPU and is more accurate than an fp16 7B at comparable memory. We further extend the kernel to additive (AQLM-style) vector codebooks: it decodes real AQLM weights, and once the codebooks are trained with beam search, the scheme beats the scalar codebook on accuracy at equal bits and equal decode speed. As a scale probe, the same pure-Rust pipeline compresses DeepSeek-R1 671B (1.34 TB bf16) to a 326 GB 4-bit artifact and decodes it coherently on a single node, streaming routed experts from disk within a 73 GB peak RAM footprint. We also report negative results, including a Hopper thread block cluster variant that came out roughly 100x slower. All code, raw numbers, and figures are public and reproducible with a single command.

1 Introduction

At batch size one, LLM text generation reads every model weight from high bandwidth memory (HBM) once per output token. The cost is dominated by byte movement, not arithmetic. On an H100 (about 3.3 TB/s, roughly 1 PFLOP/s in half precision), a 7B model in fp16 needs about 14 GB/3.3 TB/s $\approx$ 4.2 ms to read its weights per token, against about 14 μ s of compute (1.4×10^{10} FLOP at 10^{15} FLOP/s). The arithmetic units are idle more than 99% of the time. The lever for decode latency is therefore *bytes read per token*, which is exactly why weight quantization helps decode.

This report measures, rather than assumes, how much a fused low-bit kernel actually buys, and where. We contribute:

1. A fused 4-bit codebook GEMV kernel that reads one quarter of the weight bytes and decodes **2.2 to 2.34x** faster than cuBLAS fp16 on bandwidth limited GPUs (RTX 4090, A40), with numerical agreement to about 3×10^{-4} relative error.
2. A *bandwidth law*: the speedup shrinks to parity on bandwidth rich GPUs (H100), where fp16 GEMV is already near peak. The win lives on the hardware most users actually run.
3. An end-to-end realization. A naive per-layer integration *loses* (0.73x); captured as a single CUDA graph over a static-cache decode loop, the way serving engines run, the kernel reaches a **2.0x** (Llama-2 7B) to **2.45x** (13B) end-to-end decode speedup at about a third of the memory, and we show exactly why the naive path fails.
4. An extension to additive (AQLM-style) vector codebooks: a fused decode kernel that beats cuBLAS at 2 bits (**4.30x** on A40) and decodes real AQLM weights, where the trained scheme beats the scalar codebook on accuracy at equal bits and equal decode speed.
5. A fair speed vs accuracy vs memory benchmark of quantization methods, plus honest negative results from kernel variants that did not work.
6. *Speculative decoding as nearly free throughput*. Because decode is bandwidth bound, the same weight read that emits one token can verify several. We add speculative decoding [11] on top of the 4-bit runtime and validate the batched verify *lossless* on CUDA ($K=1, 2, 3$) and Metal. We measure that a small **160M** drafter *beats* a 1.1B one ($K=2$, **1.50x** vs 1.25x projected): the binding constraint is draft cost, not acceptance rate. The naive two-model wall-clock first *lost* (0.34x); profiling traced it to a register-spilling runtime- M verify kernel, and templating it on M flips the measurement to a **1.32x wall-clock win**, lossless – the integration law demonstrated and then closed at the loop level.
7. A *pure-Rust runtime that realizes all of the above*. A from-scratch runtime (CUDA and Metal behind one 24-function backend seam) runs nine architectures with *no Python at inference*, including the modern hard cases: Multi-head Latent Attention and Mixture-of-Experts (DeepSeek-V2-Lite), Gemma-2 (GeGLU, logit softcapping, four-norm residual), and Phi-4 (14B). As a scale probe, it loads and decodes **DeepSeek-R1 671B** (the full V3 architecture: low-rank query projection, 256 routed experts, sigmoid `noaux_tc` routing) on a single node, with the 4-bit expert indices left `mmap`-backed on disk and paged in on demand: 73 GB of peak host RAM against a 326 GB artifact. The compression path is torch- and transformers-free, reading `safetensors` directly, so the runtime carries no ML-framework dependency at all.

2 Background: codebook quantization and fused decode

Codebook (non uniform, k-means) weight quantization stores each weight as a small cluster index into a per output channel table:

$$W_{\text{deq}}[i, j] = \text{codebook}[\text{indices}[i, j], j], \quad Y[m, j] = \sum_i X[m, i] W_{\text{deq}}[i, j]. \quad (1)$$

Here indices are 8-bit or packed 4-bit, and the codebook holds K representative fp16 values per column ($K \in \{16, \dots, 256\}$). This is the SqueezeLLM family [2]. Dequantizing in a separate pass (write the full W , then call GEMM) wastes the bandwidth the quantization just saved. The kernel folds the lookup into the matrix product so dequantization adds no extra global memory traffic.

3 The kernel

The decode kernel computes a fused 4-bit codebook GEMV. Each thread reads a `uint32` of packed indices (eight 4-bit indices, eight output columns), so a warp reads a full 128 byte cache line over 256 columns. The per column codebook tile is staged once in local shared memory; the contraction is split across the K dimension (split-K) with an atomic float reduction. The full precision weight matrix is never written to global memory (Figure 1). Building per architecture (`sm_86` for A40, `sm_89` for RTX 40 series, `sm_90` for Hopper) the same source compiles and runs unchanged.




Figure 1: **Fused codebook decode.** The baseline dequantizes the whole weight matrix to global memory (HBM) and then calls a dense GEMM, paying an extra write and read of the full-precision W and undoing the bandwidth the quantization saved. The fused kernel decodes each 4-bit index from the per-column codebook in registers and accumulates directly into the output, so W is never materialized and only the 4-bit indices (plus the small, cached codebook) are read from HBM.

4 Experimental setup

All randomness is generated by `std::mt19937(0)` (fixed seed 0). Timings use CUDA events over 50 to 500 iterations after warmup. Every kernel is checked for numerical correctness against cuBLAS on the same random data, and we report the maximum relative error. Matrices are $IC = OC = 4096$, fp16 activations and output, decode at batch $M = 1$. The cross-GPU runs use a live cuBLAS fp16 GEMV as the baseline in the same binary.

5 Results

5.1 Cross-GPU decode and the bandwidth law

Table 1 and Figure 2 report the fused 4-bit GEMV against a live cuBLAS fp16 GEMV on three GPUs. The kernel moves one quarter of the weight bytes, so on memory bound parts it wins about 2.2 to 2.3x. On the H100 the fp16 GEMV is already near peak bandwidth, so the kernel only ties and the gain there is memory, not latency.

5.2 A40 microbenchmarks

Table 2 breaks down the A40 decode, dequantization, and prefill regimes. Decode is memory bound, so fewer index bits directly buy latency. Standalone dequantization improves 6.7x with an L2 cached gather. Prefill is compute bound: cuBLAS already runs near the Tensor Core peak, so

GPU	bandwidth class	speedup vs cuBLAS fp16	rel. err
RTX 4090 (sm_89)	~ 1.0 TB/s	2.20x	3×10^{-4}
A40 (sm_86)	~ 0.7 TB/s	2.34x	3×10^{-4}
H100 PCIe (sm_90)	~ 3.3 TB/s	0.99x (parity)	3×10^{-4}

Table 1: Fused 4-bit codebook GEMV decode speedup, 4096×4096 , $M = 1$, seed 0. The more bandwidth limited the GPU, the larger the win.




Figure 2: Decode speedup of the fused 4-bit GEMV versus cuBLAS fp16, measured per GPU. The dashed line is cuBLAS parity. The win scales inversely with available memory bandwidth.

a fused kernel can at best match it, and here reaches about $0.21 \times$ cuBLAS, ten times the naive baseline. In the compute bound regime, codebook quantization buys memory, not speed.

5.3 Speed vs accuracy vs memory of quantization methods

To place the kernel work in context, we ran a fair benchmark of three quantization methods (fp16, AWQ [3], AQLM [4]) through one harness (same HuggingFace forward path, greedy, batch 1, seed 0), reporting decode throughput, peak VRAM, and full wikitext-2 perplexity (PPL), on an H100 (Table 3). The fp16 7B PPL of 5.47¹ and AWQ 70B PPL of 3.41 match published values, which validates the harness. Comprehensive accuracy benchmarks of post-training quantization already exist and cover far more methods, bitwidths, and architectures than we do [1]; our contribution here is the orthogonal axis they do not measure, namely decode speed and the memory versus latency trade-off.

Two observations stand out. First (Figure 3), at 7B on an H100 every quantized method decodes *slower* than fp16: the bandwidth is high enough that fp16 is already fast and the dequant overhead dominates at batch 1. Quantization there is a memory play. Second (Figure 4), the story flips at

¹This 5.47 is measured with the HuggingFace forward path at sequence length 2048, median of 3 runs (the cross-method comparison table). The 5.83 reported later for the model-level ablations comes from a different harness (the custom codebook evaluation harness); both are correct for their respective protocols.

Regime	Variant	Metric	vs cuBLAS
Decode (GEMV, $M=1$)	cuBLAS fp16 (baseline)	0.0612 ms	1.00x
	uint8, $K=256$	0.0562 ms	1.09x
	uint8, $K=64$	0.0389 ms	1.57x
	4-bit, $K=16$	0.0261 ms	2.34x
Dequant (bandwidth)	naive staging	31.8 GB/s	1.0x
	L2 cached gather	213 GB/s	6.7x
Prefill (GEMM, $M=2048$)	best fused (Tensor Core)	22.4 TFLOP/s	0.21x
	cuBLAS fp16 (baseline)	~ 105 TFLOP/s	1.00x

Table 2: A40 (sm_86, CUDA 11.8), 4096×4096 , seed 0, all variants verified against cuBLAS.

Model	Method	decode (tok/s)	VRAM (GB)	wikitext-2 PPL
Llama-2 7B	fp16	43.7	15.3 ²	5.47
	AWQ 4-bit	26.8	5.7	5.60
	AQLM 2-bit	25.0	4.2	6.34
Llama-2 70B	fp16	~ 138 GB, does not fit on one 80 GB GPU		
	AWQ 4-bit	9.1	38.4	3.41
	AQLM 2-bit	9.2	20.7	4.06

Table 3: Speed, memory, and accuracy on H100, full wikitext-2 PPL, seed 0.

70B: fp16 needs about 138 GB (two or more GPUs), while AQLM 2-bit fits in 20.7 GB on a single 24 GB consumer card, and at PPL 4.06 it is markedly more accurate than an fp16 7B (PPL 5.47) at comparable memory. For a fixed VRAM budget, a 2-bit 70B dominates an fp16 7B.

6 Model-level evaluation

We quantized all 224 projection layers (q, k, v, o, gate, up, down across 32 blocks) of Llama-2 7B with the per-output-channel codebook ($K = 16$, 4 bits, per-column k-means) and evaluated end to end on an RTX 4090.

Accuracy. wikitext-2 perplexity goes from 5.83³ (fp16) to 6.34 (codebook 4-bit), a 0.51 increase. The model stays usable, but the scheme trails activation-aware AWQ (about 0.1 PPL at 4 bits), as expected for a scalar codebook trained on raw weight error without calibration or outlier handling. A simple activation-aware calibration (weighting the per-column k-means by mean activation magnitude, the core idea of AWQ) closes about a third of the gap, from 6.34 to 6.17. Going further does not help. A per-channel scale search (grid over the scale exponent) lands at 6.18, and the complete AWQ pipeline, per-channel scaling and weight clipping both selected by the real output error on cached activations, does not improve either (6.21, marginally worse). The reason is structural: AWQ’s per-channel scaling is designed for a fixed uniform quantization grid, where it redistributes resolution toward important channels. A codebook is already adaptive per column (k-means places centroids on each column’s distribution), so the scaling buys little and

³This 5.83 fp16 baseline is measured with the custom codebook evaluation harness used for these ablations, distinct from the 5.47 reported earlier under the HuggingFace forward path (seqlen 2048, median of 3); both are correct for their respective protocols.




Figure 3: Llama-2 7B on H100: decode throughput vs perplexity, bubble area equal to peak VRAM. fp16 is fastest and most accurate but memory heavy; quantization trades decode speed for memory in this regime.

the extra calibration only adds noise. Closing the gap to vector and trellis methods (AQLM, QTIP) needs a better quantizer, not AWQ’s tricks bolted onto a scalar codebook. A naive vector quantizer confirms this from the other side: quantizing pairs of weights against a single shared 256-entry codebook (iso-bit at 4 bits) is catastrophic on Llama-2 7B (perplexity diverges, with or without per-channel normalization), because one shared codebook cannot match the per-column adaptivity of the scalar scheme and lacks the incoherence processing and multi-codebook structure that make AQLM and QuIP# work. Even incoherence processing (a random orthogonal rotation on the input dimension before quantizing, the QuIP# lever, with the inverse folded back into the effective weight) moves the scalar codebook only marginally, from 6.34 to 6.29, short of the simple calibration (6.17): incoherence pays off paired with vector or lattice quantization, not on a scalar codebook alone. The served model is otherwise healthy (it generates coherent on-topic text at PPL 6.34). The accuracy frontier is a separate research program, not a quick add-on; the value of this scheme is memory and kernel speed.

Memory. Peak VRAM drops from 13.56 GB to 4.63 GB (2.9x), so the 7B model runs in under 5 GB on a consumer card.

Speed. Per GEMV the kernel is faster than dense fp16 on the large layer shapes that dominate an LLM. But a naive per-layer custom-op swap is overhead bound: end-to-end decode goes from 44 to 32 tok/s (0.73x). The per-op launch and dispatch overhead, paid 224 times per token, dominates, while the fp16 baseline uses a fused optimized path. Realizing the per-GEMV speedup at the model level needs framework integration (CUDA graphs, fused dispatch), not a per-layer swap. We verified the kernel is CUDA-graph compatible (capture is numerically correct once the kernel launches on the captured stream), but in an isolated 224-GEMV decode loop a graph adds only about 6 percent over eager, so launch overhead is not the dominant cost. Against the real cuBLAS GEMV path (F.linear), the kernel runs the 224-GEMV decode work at 172 tok/s versus 70 (2.4x), both CUDA-graphed, so the kernel genuinely beats cuBLAS at batch 1. The end-to-end 0.73x came from the naive wrapper’s per-op float32-to-float16 cast and copy, not the kernel. We then built the clean integration: the kernel writes fp16 into preallocated buffers (no per-op cast




Figure 4: Llama-2 70B: fp16 does not fit on one 80 GB GPU, quantized it fits on one, and AQLM 2-bit (21 GB) fits a single RTX 4090.

or allocation) and the decode step is captured as a single CUDA graph over a static-cache decode loop, the way serving engines run. Under this integration the codebook model decodes 123.4 versus 61.6 tok/s against fp16 on an RTX 4090 (2.0x end-to-end) at 4.73 versus 13.58 GB. So the op-level 2.4x is realized as a 2.0x end-to-end decode speedup once the per-token python overhead is removed; the arc is 0.73x (naive) to 0.85x (cast-free eager) to 2.0x (CUDA-graphed). The kernel makes a real 7B decode twice as fast at a third of the memory. The win grows with model size and bandwidth-boundedness: Llama-2 13B on an A40 decodes 49.0 versus 20.0 tok/s (2.45x) at 8.50 versus 26.17 GB (3.08x less), consistent with the bandwidth law (bigger weights read, more bandwidth-limited GPU, larger gain).

Energy. The byte savings show up as decode energy. On the RTX 4090 (batch one, 512-token generation, three runs), a hardened measurement (trapezoidal integral of sampled GPU power, fixed seed, the 59.4 W idle baseline subtracted) gives fp16 at 5.31 ± 0.13 and the 4-bit codebook (via a Python op) at 3.94 ± 0.08 J/token; net of the shared idle floor the active-decode energy is 4.03 versus 1.86 J/token, a **2.17 $\times$** reduction (the 4-bit kernel draws 112 W active versus 246 W for fp16). The standard deviation is about 2 to 3%, so the gap is well outside run-to-run noise. The pure-Rust runtime, which removes the per-op Python dispatch, lowers energy further to 2.58 J/token over the same generation.

Scale probe: DeepSeek-R1 671B. To test whether the pipeline holds at frontier scale we exported and ran DeepSeek-R1 671B, the largest open-source model, end to end. A streaming, resumable, torch-free exporter compresses the 1.34 TB bf16 checkpoint to a 326 GB 4-bit codebook artifact in about 7 hours on one A100 pod, with bounded RAM. The runtime then loads all 61 layers in 43.3 s and decodes coherently (prompt “*The capital of France is*” continues “*Paris. Paris is located in northern France and is known for its iconic landmarks such as the Eiffel Tower, Notre*”). The key mechanism is that routed-expert index tensors are never loaded: they stay `mmap`-backed on disk and the top-8 experts a token actually routes to are paged in on demand, while the small f32 codebooks stay resident. Peak host RAM is **73 GB** against the 326 GB artifact (the same load without `mmap` needs 326 GB and was killed by the container’s 250 GB limit). First-pass decode is 0.2 tok/s, entirely dominated by first-touch page faults on network-attached storage; steady-state

throughput after cache warm-up is left to future work. The point of the probe is architectural: a 671B Mixture-of-Experts runs in a from-scratch Rust runtime on one box with commodity RAM, because at batch one only the routed experts’ bytes must move. The same artifact then ran on a single consumer RTX 4090: **20.9 GB** peak VRAM (it fits a 24 GB card with margin), 0.24 tok/s steady state over 64 tokens with identical coherent output, and only 9 GB of anonymous host RAM beyond the page cache. Decode is disk-bound, not GPU-bound; the probe demonstrates reach, not serving throughput.

7 Additive vector quantization: a fast kernel that beats the scalar codebook

The scalar codebook’s accuracy ceiling is structural: per-output-channel scalar quantization cannot match vector and trellis methods (AQLM, QuIP#, QTIP) at low bits. Those methods quantize a *group* of weights jointly; AQLM reconstructs a group of D input weights as a sum of M codebook vectors, $w_g = \sum_{m=1}^M C_m[b_m]$, reaching near-fp16 quality at 2 bits. Their decode kernels are the bottleneck. We show the fused-decode idea transfers, the resulting kernel beats cuBLAS, and the scheme, once trained, beats the scalar codebook.

The kernel. For $y[o] = \sum_g \langle x_g, w_g \rangle$ with the additive reconstruction, the dot product $\langle x_g, C_m[k] \rangle$ is independent of the output channel, so a per-group lookup table $\text{LUT}[m][k] = \langle x_g, C_m[k] \rangle$ is built once in shared memory and $y[o] = \sum_g \sum_m \text{LUT}[m][b_m(o, g)]$. The kernel reads the *codes*, never the dense weight: at 2 bits ($M = 2$, $K = 256$, $D = 8$) that is 4.2 MB of codes versus 33.6 MB of fp16 weight. Against a live cuBLAS fp16 GEMV at 4096×4096 (relative error 2.3×10^{-4} , verified):

GPU	2-bit ($M = 2$)	4-bit ($M = 4$)
RTX 4090 (~ 1.0 TB/s)	1.71x	–
A40 (~ 0.7 TB/s)	4.30x	2.39x

The key optimization was vectorized 32-bit code reads (four 8-bit codes per thread), which lifted the A40 from 2.35x to 4.30x: the 1-byte reads, not synchronization, were the bottleneck. The format is exactly AQLM’s **2x8**: loading a real **Llama-2-7b-AQLM-2Bit-2x8** checkpoint, its codebooks [2, 256, 1, 8], codes [4096, 512, 2] and per-output scales [4096] map one-to-one onto the kernel, so it decodes real AQLM weights.

Accuracy, beating the scalar codebook at equal bits. The additive *structure* alone is not enough: a greedy residual fit at $M = 4$ (4 bits) only ties the scalar 4-bit codebook (PPL 6.32 vs 6.34). The accuracy comes from AQLM’s *training*. We reproduced its core (per-output scale, beam-search code assignment, least-squares codebook updates, alternating; no block fine-tuning): at $M = 4$ this reaches PPL 6.13, beating the scalar 4-bit codebook while the kernel still decodes it at 2.39x.

scheme (4-bit, Llama-2 7B)	wikitext-2 PPL
scalar codebook	6.34
greedy additive	6.32
beam-search + LSQ (this work)	6.13
fp16 (custom codebook harness)	5.83

So trained additive VQ is more accurate than the scalar codebook at the same bit rate and the same decode speed, a Pareto improvement.

Honest limits. This is one alternating round without the levers that close the remaining gap to fp16 (5.83). A naive diagonal-Hessian calibration of ours actually *regresses* (PPL 8.6), because AQLM uses the full Hessian in a sequential GPTQ-style error correction, not a diagonal MSE weighting; reproducing that, plus block fine-tuning, is a separate effort. What we establish is narrower and solid: a fused additive-VQ decode kernel that beats cuBLAS at AQLM bit rates and decodes real AQLM weights, and that the additive scheme, once trained, beats the scalar codebook at equal bits and equal speed.

8 Negative results

We keep failed variants in the repository, because they constrain the design space.

- **Hopper cluster with distributed shared memory.** Staging the codebook once in a thread block cluster and reading it remotely (distributed shared memory) came out about 100x slower than cuBLAS (1.53 ms vs 0.02 ms), though numerically correct. The codebook lookup is in the hot loop, and a remote distributed shared memory read costs far more than a local shared read. The codebook is only 8 KB, so local staging was never the bottleneck.
- **L2 resident codebook and activation staging.** Reading the codebook from L2 instead of staging it locally regressed to 0.41 to 0.60x on H100; staging the activation vector in shared memory regressed from 0.99x to about 0.90x by raising shared memory pressure and lowering occupancy. The local shared memory design remains the best, and beating cuBLAS on H100 would require a different approach (an atomic free reduction, or a TMA and wgmma path), not these tweaks.

9 Related work

Codebook quantization begins with the scalar per-output-channel scheme of SqueezeLLM [2] and extends to additive and lattice codebooks: AQLM [4], QuIP# [5], and QTIP [6] push low-bit accuracy with vector and trellis structure. On the kernel side, the relevant line is fused low-bit GEMM/GEMV. Marlin [7] and its Hopper successor Machete target uniform GPTQ [8] 4-bit; FLUTE [9] fuses lookup-table dequantization with offline weight restructuring and table duplication; and VQ-LLM [10] generates fused vector-quantization kernels around a *codebook cache* that places codebook entries across the memory hierarchy. VQ-LLM is the closest to our additive-codebook decode and, like us, finds that keeping the codebook on chip is the lever.

We differ in emphasis rather than claiming a faster kernel than these peers. First, we measure a cross-GPU *bandwidth law* against a live, same-binary cuBLAS baseline, showing the win concentrates on bandwidth-limited consumer parts and vanishes to parity on an H100. Second, we foreground the end-to-end integration lesson that kernel papers usually omit: a naive per-layer swap loses (0.73 $\times$), and recovering the microbenchmark win requires capturing the decode step as a single CUDA graph. Third, we add a Python-free Rust runtime with measured per-token energy. Fourth, on accuracy we show a beam-trained additive codebook beating the scalar codebook at equal bits and equal decode speed.

10 Baselines against peer quantized kernels

Our primary speed baseline is deliberately a live, same-binary cuBLAS fp16 GEMV: it is the fair systems reference for the batch-one, bandwidth-bound decode regime, runs on the identical shapes

and fixed seed, and cannot be tuned in our favor. Here we state explicitly how the kernel relates to the two peer families a reviewer will ask about, AQLM and Marlin, separating what is measured from what is not.

AQLM: matched accuracy by construction. Our additive-VQ kernel does not approximate AQLM, it decodes the *same* weights. It loads a real Llama-2-7b-AQLM-2Bit-2x8 checkpoint whose codebooks [2, 256, 1, 8], codes [4096, 512, 2] and per-output scales map one-to-one onto the kernel, so accuracy is *identical* to AQLM at 2 bits by construction and only decode speed is in question. Against same-binary cuBLAS fp16, the kernel runs these real AQLM weights at $4.30\times$ (A40, 2-bit) and $2.39\times$ (4-bit). The remaining comparison, against AQLM’s *own* kernel (`aqlm::code2x8_matmat`) on the identical Llama-2 7B attention shape, is wired op-level in `bench/avq_oplevel.py`; completing that matched-bits, matched-accuracy head-to-head is our highest-priority open item, and we report it here as unfinished rather than claim it as a win.

Marlin and Machete: a different accuracy regime. Marlin [7] and its Hopper successor Machete are fused W4A16 kernels for *uniform* GPTQ [8] 4-bit, reaching near-ideal $\sim 4\times$ dense-GEMM speedups on Ampere as reported by their authors. They are not a matched-accuracy competitor to the low-bit vector regime where codebook and additive-VQ schemes live: uniform scalar 4-bit and 2-bit AQLM sit at different points on the accuracy curve, so a raw speed ratio between them would compare kernels solving different problems. At uniform 4-bit our scalar codebook is the comparable accuracy point, and there Marlin’s tuned dense GEMM is the kernel to beat on *batched* throughput, which we do not claim: our wins are in the batch-one GEMV regime, and a batched GEMM path is an open problem. We therefore position against Marlin by regime, not by an unmeasured number.

Peer kernel	Quant regime	Accuracy vs ours	What we measure
cuBLAS fp16	dense fp16	exact fp16 reference	$2.0\text{--}4.30\times$ decode, same binary
AQLM official	2-bit vector (2x8)	identical weights	decode the same checkpoint at $4.30\times$ vs fp16; op-level vs their kernel in progress
Marlin / Machete	uniform GPTQ 4-bit	different regime	not measured; positioned by design (batched GEMM, open)
FLUTE, VQ-LLM	LUT / vector	different regime	not measured; closest peers, discussed in related work

In short: against AQLM the accuracy is fixed and we already decode its real weights faster than fp16, with the kernel-vs-kernel timing the one number still owed; against Marlin the honest statement is that it targets a different bit-accuracy operating point and a batched regime we do not yet contest.

11 Speculative decoding: harvesting the idle compute

The bandwidth law that motivates the fused kernel has a second consequence. At batch one the arithmetic units sit nearly idle while the weights stream from memory; the same weight read that produces one token could verify several. Speculative decoding [11] exploits exactly this: a small drafter proposes K tokens, and the target verifies all $K+1$ in a single forward that reads its weights once. Because decode is bandwidth bound, that verify is close to free.

Verifying $K+1$ positions in one forward requires a small- M ($M = K+1$) batched variant of every decode operation: the fused codebook GEMM, RMSNorm, RoPE, causal attention, and

the KV-cache append. We implement these ($M \leq 4$) in both the CUDA and Metal backends and validate that the batched forward and the resulting speculative loop are *lossless*: the emitted sequence is token-for-token identical to plain greedy decode, for $K = 1, 2, 3$ and for both cooperative and adversarial drafters. A rejected draft leaves stale KV rows that the next step overwrites, so no explicit rollback is needed.

We measure the acceptance rate α (the fraction of positions where the drafter’s greedy token matches the target’s) between a 4-bit Llama-2-7B target and two 4-bit drafters on an RTX 4090, and project the speedup with the standard estimate $S(K) = \frac{1-\alpha^{K+1}}{1-\alpha} / (1 + K t_{\text{draft}}/t_{\text{target}})$, where the verify counts as one target forward. TinyLlama-1.1B gives $\alpha = 0.73$ and a best $1.25\times$ at $K = 2$; the much smaller **llama-160m** gives a lower $\alpha = 0.65$ but, being far cheaper per draft, a better $1.50\times$ **at** $K = 2$. The smaller drafter wins: draft cost, not raw acceptance, is the binding constraint, and $K = 2$ (verify three tokens per target forward) is optimal for both. This is lossless throughput, orthogonal to and compounding with the quantization memory and energy savings.

Wall-clock: lost, diagnosed, fixed, won. We then built the real two-model harness (the drafter is a second model decoding incrementally in its own KV cache, with longest-common-prefix row reuse) and measured it end to end on the same 4090. The naive loop first *lost* $3\times$ ($0.34\times$ plain decode at $\alpha = 0.91$ effective acceptance, lossless at every K) — the integration law of the model-level evaluation repeating at the loop level. Profiling isolated the cause: the batched-verify kernel took M at runtime, which defeats unrolling and spills the accumulator array to local memory; the $M \leq 4$ verify cost $5.5\times$ ($M=2$) to $10.7\times$ ($M=4$) of one $M=1$ GEMV. Making M a compile-time template parameter keeps the accumulators in registers and restores the bandwidth argument: the $M=2$ verify then costs $1.11\times$ an $M=1$ forward, and speculative decoding measures $1.32\times$ **wall-clock at** $K=2$ (164 vs 125 tok/s, token-for-token lossless) on the Llama-2-7B + 160m pair, against the $1.50\times$ projected ceiling. Large-vocabulary pairs (Qwen2.5, 152k; Llama-3.1, 128k) currently sit at $0.9\times$ parity: every drafter forward reads back the full logits for a host argmax (600 KB per token at 152k vocab), which a device-side argmax removes — the remaining step, together with graph capture of the whole speculative round.

12 Limitations and scope

This is a kernel and benchmark study, not a new quantization method. The scheme is a scalar per output channel codebook (SqueezeLLM family), which trails vector and trellis methods (QTIP [6], AQLM [4], QuIP# [5]) on low-bit accuracy. The accuracy harness uses the HuggingFace forward path, not a tuned serving engine, so absolute throughput is not a serving claim. Our headline kernel baseline is cuBLAS fp16; the comparison against specialized peers (Marlin [7], Machete, FLUTE [9], VQ-LLM [10]) is only partial, as detailed in the baselines section: we decode AQLM’s real 2-bit weights faster than fp16, but the op-level timing against AQLM’s own kernel is unfinished and we make no matched-accuracy speed claim against Marlin. Prefill stays at about $0.21\times$ cuBLAS, compute bound as expected. Numbers were taken at a single 4096×4096 shape and will differ at other shapes. The microbenchmarks also use random indices and codebooks; real k -means indices and weight distributions may shift cache and bank-conflict behavior, which we have not isolated from the synthetic case.

13 Reproducibility and code

All code is public as **Trapetum**, a single repository that builds and runs with one command and captures GPU, driver, CUDA version, and seed in its outputs. The CUDA kernels are in **kernels/**,

the model-level scripts in `model/`, the quantization benchmark in `bench/`, the pure-Rust runtime in `runtime/`, and this paper in `paper/`.

<https://github.com/neuralboot/trapetum>

<https://neuralboot.com/trapetum>

14 Conclusion

At batch one, decode is bandwidth bound and the arithmetic units are nearly idle. A fused 4-bit codebook GEMV that reads one quarter of the weight bytes turns that idle headroom into a 2.2 to 2.34x decode speedup on the GPUs most people run, while tying cuBLAS on bandwidth rich datacenter parts. Captured as a CUDA graph the way serving engines run, this becomes a 2.0x (7B) to 2.45x (13B) end-to-end decode speedup on real models, where a naive per-layer integration instead loses. Extending the kernel to additive vector codebooks, the trained scheme beats the scalar codebook at equal bits and equal speed, a step toward fast and accurate low-bit decode. The complementary quantization benchmark shows the same theme from the model side: low-bit quantization is a memory lever first, and on large models it is what makes the model runnable at all, a 2-bit 70B on a single consumer card. The honest negative results map where the easy speedups are not.

References

- [1] J. Zhao, M. Wang, M. Zhang, Y. Shang, X. Liu, Y. Wang, M. Zhang, and L. Nie. Benchmarking Post-Training Quantization in LLMs: Comprehensive Taxonomy, Unified Evaluation, and Comparative Analysis. arXiv:2502.13178, 2025.
- [2] S. Kim, C. Hooper, A. Gholami, et al. SqueezeLLM: Dense-and-Sparse Quantization. arXiv:2306.07629, 2023.
- [3] J. Lin, J. Tang, H. Tang, S. Yang, et al. AWQ: Activation-aware Weight Quantization for LLM Compression and Acceleration. Proc. MLSys, 2024. arXiv:2306.00978.
- [4] V. Egiazarian, A. Panferov, D. Kuznedelev, E. Frantar, A. Babenko, and D. Alistarh. Extreme Compression of Large Language Models via Additive Quantization. Proc. ICML, 2024. arXiv:2401.06118.
- [5] A. Tseng, J. Chee, Q. Sun, V. Kuleshov, and C. De Sa. QuIP#: Even Better LLM Quantization with Hadamard Incoherence and Lattice Codebooks. Proc. ICML, 2024. arXiv:2402.04396.
- [6] A. Tseng, Q. Sun, D. Hou, and C. De Sa. QTIP: Quantization with Trellises and Incoherence Processing. Proc. NeurIPS, 2024. arXiv:2406.11235.
- [7] E. Frantar, R. L. Castro, J. Chen, T. Hoefer, and D. Alistarh. MARLIN: Mixed-Precision Auto-Regressive Parallel Inference on Large Language Models. arXiv:2408.11743, 2024.
- [8] E. Frantar, S. Ashkboos, T. Hoefer, and D. Alistarh. GPTQ: Accurate Post-Training Quantization for Generative Pre-trained Transformers. Proc. ICLR, 2023. arXiv:2210.17323.
- [9] H. Guo et al. Fast Matrix Multiplications for Lookup Table-Quantized LLMs. Findings of EMNLP, 2024. arXiv:2407.10960.

- [10] VQ-LLM: High-Performance Code Generation for Vector Quantization Augmented LLM Inference. arXiv:2503.02236, 2025.
- [11] Y. Leviathan, M. Kalman, and Y. Matias. Fast Inference from Transformers via Speculative Decoding. ICML, 2023. arXiv:2211.17192.